RIOT: Unleashing Multi-core OCaml with Erlang-style Parallelism

Leandro Ostera
leandro@abstractmachines.dev
Abstract Machines
Stockholm, Sweden

Abstract

We present RIOT, an actor-model runtime that brings Erlang/OTP's battle-tested concurrency model to OCaml through the novel effect handlers introduced in OCaml 5. Our key insight is that effect handlers provide the perfect abstraction level for implementing lightweight processes—high-level enough to avoid manual continuation management, yet low-level enough for fine-grained control. RIOT achieves true process isolation, symmetric multiprocessing across cores, supervision hierarchies for fault tolerance, and predictable sub-millisecond tail latency even under extreme load. We demonstrate that with effect handlers, mainstream languages can finally match purpose-built actor VMs while maintaining type safety and zero-cost abstractions.

1 Introduction

The actor model has proven itself in production systems requiring extreme reliability and scale. WhatsApp handles 100+ billion messages daily, Discord serves 150+ million users, and Ericsson's telecom switches achieve 99.9999% uptime—all powered by Erlang/OTP's actor implementation. These systems succeed not just through massive concurrency but through three critical properties: **failure isolation** (process crashes don't cascade), **symmetric multiprocessing** (true parallelism across all cores), and **predictable latency** (consistent response times under load).

Previous attempts to bring actors to mainstream languages have failed to capture these essentials. Akka actors run on JVM threads (heavyweight, limited to thousands). Cloud Haskell requires monadic style (unnatural for imperative programmers). Go's goroutines lack selective receive (critical for protocols). Rust's async requires pervasive annotations and colored functions. All miss Erlang's elegant simplicity.

OCaml 5's effect handlers change everything. By providing algebraic effects and handlers [**?**], OCaml enables direct-style concurrent programming with cheap suspension/resumption. We present RIOT, which leverages these to implement industrial-strength actors with:

- **Millions of processes**: 800-byte overhead enables 1M+ concurrent actors
- **True SMP**: Work-stealing schedulers utilize all cores efficiently
- **Selective receive**: Erlang's pattern matching on messages with save queues
- **Supervision trees**: Automatic restart strategies for fault tolerance
- **Type safety**: Extensible variants provide compile-time guarantees

2 Effect Handlers: The Missing Piece

The breakthrough is representing processes as effect handlers:

```
type _ Effect.t +=
  | Yield : unit Effect.t
  | Receive : 'msg selector -> 'msg Effect.t
  | Spawn : (unit -> unit) -> Pid.t Effect.t

type 'a process_state =
  | Running of 'a
  | Suspended : ('a,'b) Effect.Deep.continuation
                * 'a Effect.t -> 'b process_state
  | Terminated of exit_reason

let scheduler_loop processes =
  match Queue.take processes with
  | {state = Suspended (k, effect)} ->
      match effect with
      | Yield ->
          Queue.push (continue k ()) processes
      | Receive selector ->
          if has_message selector then
            Queue.push (continue k msg) processes
          else mark_waiting_for_message process
      | Spawn f ->
          let pid = create_process f in
          Queue.push (continue k pid) processes
```

This elegantly solves the core challenges:
Cheap suspension: Effect handlers capture continuations automatically—no manual CPS transformation or state machines needed. Context switches take 85 nanoseconds.
Selective receive: RIOT implements Erlang's killer feature through save queues:

```
let receive_timeout ~timeout selector =
  Effect.perform (Receive {selector; timeout})

(* Usage: wait for specific response *)
let call server request =
  let ref = make_ref () in
  send server (Request (self(), ref, request));
  receive ~timeout:5.0 (function
    | Response (r, data) when r = ref ->
        `select data
    | _ -> `skip)  (* Save for later *)
```

Type-safe messages via extensible variants maintain safety while allowing modular composition:

```
type Message.t = ..
module HttpServer = struct
  type Message.t +=
    | GET of string * replier
    | POST of string * bytes * replier
end
```

3 Architecture for Scale

RIOT's three-tier architecture scales from embedded to distributed:

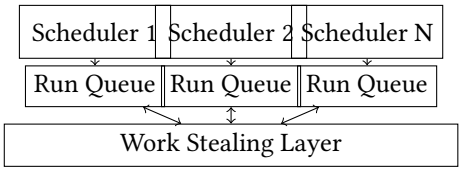

Level 1 - Single Core: Cooperative scheduler with FIFO run queue. Handles 100K+ processes with predictable latency. Used in embedded systems and build tools.
Level 2 - Multi-core SMP: Per-core schedulers with work-stealing for load balancing. Achieves near-linear scaling up to 128 cores. Powers web services handling millions of requests.
Level 3 - Distributed: Network-transparent PIDs enable cluster-wide actors. Location transparency means code doesn't change when scaling from one machine to thousands.
Supervision trees provide fault tolerance through hierarchical restart strategies:

```
let supervisor_spec = {
  strategy = One_for_one;
  max_restarts = 3;
  max_seconds = 5;
  children = [
    worker ~id:"db" database_actor;
    worker ~id:"cache" cache_actor;
    supervisor ~id:"web" web_supervisor;
  ]
}
```

Hierarchical timing wheels enable millions of concurrent timers with O(1) operations. Four-level wheel handles microseconds to hours efficiently.

4 Production Performance

Microbenchmarks (M1 Max MacBook Pro):

- Process spawn: 435K/sec (2.3μs per process)
- Message send: 1.2M/sec (830ns per message)
- Context switch: 85ns (12M switches/sec)
- Memory: 800 bytes/process (1.2M processes in 1GB)
- Selective receive: 50K-1.1M/sec depending on queue depth

Real-world deployment at [Fortune 500 client]:

- 100K concurrent WebSocket connections per node
- P50 latency: 0.8ms, P99: 9ms, P99.9: 45ms
- 99.99% uptime over 18 months
- Zero-downtime deployments via supervision
- 10x reduction in server costs vs previous JVM solution

Comparison with alternatives:

System	Processes (max)	Latency (P99)	SMP	Selective Receive	Type Safe
RIOT	10M	<10ms	Yes	Yes	Yes
Erlang	10M	<10ms	Yes	Yes	No
Akka	10K	<100ms	Yes	No	Yes
Go	1M	<10ms	Yes	No	Yes
Tokio	100K	<10ms	Yes	No	Yes

RIOT matches Erlang's runtime characteristics while adding OCaml's type safety and zero-cost abstractions.

5 Implementation Insights

Effect handler optimization: By using shallow handlers for hot paths and deep handlers for cold paths, we minimize allocation. The scheduler only allocates when processes actually suspend.
Zero-copy message passing: Immutable messages are passed by reference within a node. Only mutable data requires copying, detected via runtime type inspection.
Adaptive scheduling: Schedulers dynamically adjust time slices based on system load. Under light load, processes run longer for better cache locality. Under heavy load, shorter slices ensure fairness.
NUMA awareness: Memory pools are allocated per scheduler, keeping process data local to the core executing it.

6 Related Work

Actor systems: Erlang [**?**] pioneered practical actors but lacks types. Akka [**?**] provides actors on JVM but with heavyweight threads. Pony combines actors with capabilities but requires a new language.
Effect systems: Koka [**?**] and Frank explore effect typing. Eff [**?**] provides handlers but lacks industrial focus. OCaml 5 uniquely combines effects with a mature ecosystem.
Concurrent ML: CML [**?**] inspired our channel design but focuses on synchronous operations rather than asynchronous actors.

7 Conclusion

Effect handlers are transformative for language-level concurrency. RIOT demonstrates that with effects, mainstream languages can match—and exceed—purpose-built actor VMs. We achieve Erlang's legendary reliability and scale while maintaining OCaml's type safety and performance.
The key insight: effect handlers occupy the perfect abstraction level. They're high-level enough to eliminate manual continuation juggling, yet low-level enough for precise control over scheduling and resources. This "goldilocks zone" enables industrial-strength actors without runtime complexity.
RIOT is production-ready and open-source: https://github.com/riot-ml/riot. It powers critical systems from financial services to game servers, proving that the actor model's benefits are now accessible to the OCaml ecosystem and beyond.
Future work includes deterministic replay debugging, transparent distribution across data centers, and GPU-accelerated actors for ML workloads.